

MDSC-102 (P) End Sem Assessment

Name: Vidyasager G R

Roll No: 23914

Assessment Tasks: Take a Dataset, and perform the following operations.

1. Data Cleaning.
2. Summary Statistics (Mean, Mean, Mode, etc)
3. Inferences through Plots.
4. Analyze Features and their Correlation.
5. Take one or more continuous features and check their distribution.
6. Take a parameter of one feature and estimate its confidence interval (for example., Mean of a Feature.)
7. Present an hypothesis for the same parameter and test it using statistical tests that fits it the most.

Loading Data

First we make a dataframe from the dataset using Pandas. The dataset that I have taken for this task is the Manga Database which is scraped from a website called the 'MyAnimeList'. 'MyAnimeList' holds a huge database of both Anime and Manga (Japanese Books / Comics). It is like iMDB for Anime and Manga.

NOTE: This dataset was scraped and uploaded in Kaggle by a third-person.

`pd.read_csv()` is used to read the data into a Pandas Dataframe.

```
df = pd.read_csv('/kaggle/input/myanimelist/manga.csv')
df.head()
```

manga_id	title	type	score	scored by	status	volumes	chapters	start_date	end_date	members	favorites	sfw	approved	created_at_before	updated_at	real_start_date	real_end_date	genres	themes	demographics	authors	serializations	
0	2	Berserk	manga	9.47	319696	currently_publishing	NaN	NaN	1989-08-25	NaN	643969	119470	True	True	2007-07-17 20:14:45+00:00	2023-04-01 00:19:31+00:00	1989-08-25	NaN	['Action', 'Adventure', 'Award Winning', 'Dram...']	['Gore', 'Military', 'Mythology', 'Psychologic...']	['Seinen']	[{'id': 1868, 'first_name': 'Kentaro', 'last_...'}]	['Young Adult']
1	13	One Piece	manga	9.22	355375	currently_publishing	NaN	NaN	1997-07-22	NaN	579557	111462	True	True	2007-07-17 20:14:45+00:00	2023-06-24 12:39:48+00:00	1997-07-22	NaN	['Action', 'Adventure', 'Fantasy']	[]	['Shounen']	[{'id': 1881, 'first_name': 'Eiichiro', 'last_...'}]	['Shounen Jump (Weekly)']
2	1706	Jōjo no Kimyō na Bōken Part 7: Steel Ball Run	manga	9.30	151433	finished	24.0	96.0	2004-01-19	2011-04-19	248511	41713	True	True	2007-10-07 08:14:20+00:00	2023-04-02 18:07:42+00:00	2004-01-19	2011-04-19	['Action', 'Adventure', 'Mystery', 'Supernatur...']	['Historical']	['Seinen', 'Shounen']	[{'id': 2619, 'first_name': 'Hirohiko', 'last_...'}]	['Ultra Jump']
3	4632	Oyasumi Punpun	manga	9.02	168459	finished	13.0	147.0	2007-03-15	2013-11-02	413897	49361	True	True	2008-02-03 15:54:30+00:00	2023-04-02 18:09:11+00:00	2007-03-15	2013-11-02	['Drama', 'Slice of Life']	['Psychological']	['Seinen']	[{'id': 2836, 'first_name': 'Ino', 'last_name': ...}]	['Big Comic Spirits']
4	25	Fullmetal Alchemist	manga	9.03	153151	finished	27.0	116.0	2001-07-12	2010-09-11	284027	29634	True	True	2007-07-17 20:14:45+00:00	2023-05-27 17:21:27+00:00	2001-07-12	2010-09-11	['Action', 'Adventure', 'Award Winning', 'Dram...']	['Military']	['Shounen']	[{'id': 1874, 'first_name': 'Hiromu', 'last_na...'}]	['Shounen Gangan']

A Brief Description of the Dataset and its Features:

manga_id - MyAnimeList Manga ID.
title - Title (rōmaji or english).
type - Manga media type.
score - MAL Weighted Score.
scored_by - Number of users who scored this manga.
status - Publishing status.
volumes - Number of Volumes.
chapters - Number of Chapters.
start_date - Publishing Start Date.
end_date - Publishing End Date.
members - Number of users with this manga in their list.
favorites - Number of users who favorited this manga.
sfw - Whether it's Safe For Work or it's R18+.
approved - Whether it's approved or is yet 'pending approval'.
created_at_before - MAL Manga entry created before this date.
updated_at - MAL Manga entry latest update date.
real_start_date - Some entries only have year or year-month.
real_end_date - Some entries only have year or year-month.
genres - List of Genres.
themes - List of Themes.
demographics - List of Demographics.
authors - List of Authors (id, first and last name, role).
serializations - List of Magazines where it's been published.
synopsis - Manga description.
background - Manga additional information.
main_picture - MAL main picture url.
url - MyAnimeList url.
title_english - Title in English.
title_japanese - Title in Japanese.
title_synonyms - List of title synonyms.

From this, I could understand what each feature in the dataset meant and so I could decide which columns/features to drop. Dropping unnecessary features will make analysis a lot easier. So, I decided to drop the following features: '**manga_id**', '**approved**', '**synopsis**', '**background**', '**genres**', '**themes**', '**demographics**', '**authors**', '**serializations**', '**main_picture**', '**url**', '**created_at_before**', '**updated_at**', '**title_english**', '**title_japanese**', '**title_synonyms**'.

This was done using the **drop()** function provided by Pandas where a list of rows (indices) /columns to be dropped are passed as parameters to the function along with another parameter '**axis**' that specifies whether the function should search entries to be removed by columns or by rows. (**axis=1 is columns and axis=0 is rows**)

```
df.drop(['manga_id', 'approved', 'synopsis', 'background', 'genres', 'themes', 'demographics', 'authors', 'serializations',
        'main_picture', 'url', 'created_at_before', 'updated_at', 'title_english', 'title_japanese', 'title_synonyms'], axis=1, inplace=True)
df
```

	title	type	score	scored_by	status	volumes	chapters	start_date	end_date	members	favorites	sfw	real_start_date	real_end_date
0	Berserk	manga	9.47	319696	currently_publishing	NaN	NaN	1989-08-25	NaN	643969	119470	True	1989-08-25	NaN
1	One Piece	manga	9.22	355375	currently_publishing	NaN	NaN	1997-07-22	NaN	579557	111462	True	1997-07-22	NaN
2	Jolo no Kimyou na Bouken Part 7: Steel Ball Run	manga	9.30	151433	finished	24.0	96.0	2004-01-19	2011-04-19	248511	41713	True	2004-01-19	2011-04-19
3	Oyasumi Punpun	manga	9.02	168459	finished	13.0	147.0	2007-03-15	2013-11-02	413897	49361	True	2007-03-15	2013-11-02
4	Fullmetal Alchemist	manga	9.03	153151	finished	27.0	116.0	2001-07-12	2010-09-11	284027	29634	True	2001-07-12	2010-09-11
...
64828	Raise wa Kimi no Mono	manga	NaN	1	finished	1.0	8.0	2018-08-30	2019-04-10	4	0	True	2018-08-30	2019-04-10
64829	Ore to Basil no Automata	manga	NaN	0	finished	2.0	19.0	2013-01-19	2014-04-19	4	0	True	2013-01-19	2014-04-19
64830	Happy Candy Virus	manga	NaN	0	finished	4.0	54.0	2021-08-02	2023-03-13	4	0	True	2021-08-02	2023-03-13
64831	Suibu Yametai Hagino-san	one_shot	NaN	0	finished	NaN	1.0	2022-03-03	2022-03-03	4	0	True	2022-03-03	2022-03-03
64832	Kidou Senshi Gundam F91 Prequel	manga	NaN	0	currently_publishing	NaN	NaN	2020-01-24	NaN	3	0	True	2020-01-24	NaN

These features were dropped as there was no relevance with the subject at hand. So, now the subset of features that I have to work with are: **‘title’, ‘type’, ‘score’, ‘scored_by’, ‘status’, ‘volumes’, ‘status’, ‘volumes’, ‘chapters’, ‘start_date’, ‘end_date’, ‘members’, ‘favorites’, ‘sfw’, ‘real_start_date’, ‘real_end_date’**.

So now, the data frame went from having 30 columns to 14 columns.

Data Cleaning

The dataset was checked for unexpected values and missing values. Although there were no unexpected values, there were a lot of missing values and I checked how many values were missing in each column using a simple piece of code.

```
df.isnull().sum()
```

```
title      0
type       0
score     40197
scored_by  0
status     0
volumes    16622
chapters   18791
start_date 1883
end_date   12901
members    0
favorites  0
sfw        0
real_start_date 1883
real_end_date 12901
dtype: int64
```

The dataset was pretty clean as it is and did not require a lot of cleaning or pre-processing from my end. Although, we could still see that there were a lot of missing values especially more than 50% of the entries were missing for the feature ‘score’. So, I dropped all the rows that had missing values for the feature ‘score’.

```
indexScore = df[df['score'].isnull()].index
indexScore
```

```
Index([24636, 24637, 24638, 24639, 24640, 24641, 24642, 24643, 24644, 24645,
      ...,
      64823, 64824, 64825, 64826, 64827, 64828, 64829, 64830, 64831, 64832],
      dtype='int64', length=40197)
```

```
df.drop(indexScore, inplace=True)
df
```

	title	type	score	scored_by	status	volumes	chapters	start_date	end_date	members	favorites	sfw	real_start_date	real_end_date
0	Berserk	manga	9.47	319696	currently_publishing	NaN	NaN	1989-08-25	NaN	643969	119470	True	1989-08-25	NaN
1	One Piece	manga	9.22	355375	currently_publishing	NaN	NaN	1997-07-22	NaN	579557	111462	True	1997-07-22	NaN
2	Jolo no Kimyou na Bouken Part 7: Steel Ball Run	manga	9.30	151433	finished	24.0	96.0	2004-01-19	2011-04-19	248511	41713	True	2004-01-19	2011-04-19
3	Oyasumi Punpun	manga	9.02	168459	finished	13.0	147.0	2007-03-15	2013-11-02	413897	49361	True	2007-03-15	2013-11-02
4	Fullmetal Alchemist	manga	9.03	153151	finished	27.0	116.0	2001-07-12	2010-09-11	284027	29634	True	2001-07-12	2010-09-11
...
24631	Fukushuu Josou Rapel! Namaiki na Kawaii Furyou...	doujinshi	5.96	102	finished	1.0	NaN	2013-01-19	2013-01-19	311	1	False	2013-01-19	2013-01-19
24632	Houkago	one_shot	5.82	103	finished	NaN	1.0	2010-01-01	2010-01-01	194	0	True	2010	2010
24633	Home Stay x Steady	one_shot	5.84	100	finished	NaN	1.0	NaN	NaN	191	0	True	NaN	NaN
24634	Oni Shimin	one_shot	5.71	101	finished	NaN	1.0	2015-05-15	2015-05-15	166	0	True	2015-05-15	2015-05-15
24635	Nichijou Sahanji no Musuko	one_shot	5.64	101	finished	NaN	1.0	1982-05-24	1982-05-24	185	0	True	1982-05-24	1982-05-24

24636 rows x 14 columns

The other features were not imputed or cleaned in someway because it actually made sense for them to be missing values, for example. The currently publishing mangas will definitely not have an end_date. The Volumes and Chapters would obviously not be shown for the currently publishing mangas. So, those were not worked with.

I although worked with the date columns by first converting them into date-time format and then just extracting the years and keeping them in their respective columns.

```
for col in ['start_date', 'end_date', 'real_start_date', 'real_end_date']:
    df[col] = pd.to_datetime(df[col], format='mixed')

# for col in ['volumes', 'chapters']:
#     df[col] = df[col].astype('Int64')
```

```
df.head()
```

	title	type	score	scored_by	status	volumes	chapters	start_date	end_date	members	favorites	sfw	real_start_date	real_end_date
0	Berserk	manga	9.47	319696	currently_publishing	NaN	NaN	1989-08-25	NaT	643969	119470	True	1989-08-25	NaT
1	One Piece	manga	9.22	355375	currently_publishing	NaN	NaN	1997-07-22	NaT	579557	111462	True	1997-07-22	NaT
2	JoJo no Kimyou na Bouken Part 7: Steel Ball Run	manga	9.30	151433	finished	24.0	96.0	2004-01-19	2011-04-19	248511	41713	True	2004-01-19	2011-04-19
3	Oyasumi Punpun	manga	9.02	168459	finished	13.0	147.0	2007-03-15	2013-11-02	413897	49361	True	2007-03-15	2013-11-02
4	Fullmetal Alchemist	manga	9.03	153151	finished	27.0	116.0	2001-07-12	2010-09-11	284027	29634	True	2001-07-12	2010-09-11

```
df['start_date'] = df['start_date'].dt.year.astype('Int64')
df['end_date'] = df['end_date'].dt.year.astype('Int64')
```

```
df.head()
```

	title	type	score	scored_by	status	volumes	chapters	start_date	end_date	members	favorites	sfw	real_start_date	real_end_date
0	Berserk	manga	9.47	319696	currently_publishing	NaN	NaN	1989	<NA>	643969	119470	True	1989-08-25	NaT
1	One Piece	manga	9.22	355375	currently_publishing	NaN	NaN	1997	<NA>	579557	111462	True	1997-07-22	NaT
2	JoJo no Kimyou na Bouken Part 7: Steel Ball Run	manga	9.30	151433	finished	24.0	96.0	2004	2011	248511	41713	True	2004-01-19	2011-04-19
3	Oyasumi Punpun	manga	9.02	168459	finished	13.0	147.0	2007	2013	413897	49361	True	2007-03-15	2013-11-02
4	Fullmetal Alchemist	manga	9.03	153151	finished	27.0	116.0	2001	2010	284027	29634	True	2001-07-12	2010-09-11

Now, that the data cleaning is done, I now check the data and try to summarize.

Summary Statistics

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24636 entries, 0 to 24635
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   title                  24636 non-null  object 
1   type                   24636 non-null  object 
2   score                  24636 non-null  float64
3   scored_by              24636 non-null  int64  
4   status                 24636 non-null  object 
5   volumes                18322 non-null  float64
6   chapters               19939 non-null  float64
7   start_date             23765 non-null  Int64  
8   end_date               19599 non-null  Int64  
9   members                24636 non-null  int64  
10  favorites              24636 non-null  int64  
11  sfw                    24636 non-null  bool   
12  real_start_date        23765 non-null  datetime64[ns]
13  real_end_date          19599 non-null  datetime64[ns]
dtypes: Int64(2), bool(1), datetime64[ns](2), float64(3), int64(3), object(3)
memory usage: 2.5+ MB
```

The data frame has 24636 rows/entries and 14 columns/features after cleaning that data. The **info()** function returns a few details about the data frame such as the amount of non-null values in each column and data type of the data stored in each column.

The **describe()** function gives us the descriptive statistics of the data such as mean, mode, median, etc and it is really useful for our purpose of making inferences.

```
df.describe()
```

	score	scored_by	volumes	chapters	start_date	end_date	members	favorites
count	24636.000000	24636.000000	18322.000000	19939.000000	23765.0	19599.0	24636.000000	24636.000000
mean	6.945179	1769.555691	3.921242	26.740358	2010.219903	2011.090719	4427.903190	143.932619
min	2.440000	100.000000	1.000000	1.000000	1946.0	1947.0	166.000000	0.000000
25%	6.620000	203.000000	1.000000	5.000000	2006.0	2007.0	610.000000	1.000000
50%	6.910000	418.000000	2.000000	9.000000	2011.0	2012.0	1191.000000	4.000000
75%	7.250000	1060.000000	4.000000	26.000000	2016.0	2017.0	2896.000000	19.000000
max	9.470000	394362.000000	200.000000	6477.000000	2023.0	2023.0	643969.000000	119470.000000
std	0.520170	8839.552408	6.005744	70.176909	8.462833	8.097007	17656.467096	1711.092022


These values will be used later for applications such as Construction of Confidence Intervals and Testing Hypothesis.

Visualizing Continuous Features


There are totally 6 continuous features in the dataset, namely 'score', 'scored_by', 'volumes', 'chapters', 'members' and 'favorites'. I used Plotly, a plotting library in python to plot the graphs.

```
columns = [['score', 'scored_by'], ['volumes', 'chapters'], ['members', 'favorites']]
fig = make_subplots(rows=3, cols=2)
for i in range(3):
    for j in range(2):
        fig.add_trace(go.Histogram(x=df[columns[i][j]], name=columns[i][j]), row=i+1, col=j+1)
fig.update_layout(height=1600, width=1600, title_text="Plotting Distributions of each Numeric Feature")
```


Plotting Distributions of each Numeric Feature



The 'score' feature was the one that made sense when plotted. I could see the bell shape in the graph, hence I could atleast just by looking, say that the data in the feature follow a normal distribution.



Plotting it using KDE helped a lot. Now the only step left was to plot a few more different types of plots on the same feature to concretize the notion that it followed a normal distribution. Checking the skewness also helped.




I plotted the 'score' feature with boxplot and a violinplot using seaborn, a yet another python library used for plotting. The boxplot and violinplot definitely helped in being sure that the feature followed a normal distribution, one additional aspect discovered was the amount of outliers that were there in the feature was a lot.


The skewness could be checked using the `skew()` function that Pandas offers, and the skewness value that I got was **0.18206351648076116**. But 0.18 felt a little bit higher for a feature that was normal, so I decided to normalize that feature even further.

Before normalizing, I tried plotting and finding the correlation between the 'score' feature and other features starting with the heatmap.

```
plt.figure(figsize=(8,6))
sns.heatmap(df.corr(numeric_only=True), annot=True)
plt.show()
```




This might be the most obvious inference that can be made. We could see from the correlation heatmap that the features 'volumes' and 'chapters' are positively correlated and to a good extent which makes sense because it is inherent that if the number of volumes are high, the book should have more chapters. A lot of the other correlation coefficients also only state the obvious for this dataset. The scatter plot shown below shows the positive correlation between volumes and chapters. All of the upcoming plots will be split into two mostly on feature 'sfw'.




There are a few more obvious inferences we can do about correlations between features that makes a lot of sense, like the number of 'members' who have added that book to their library and the number of people who have 'favourited' them; 'scored_by' and 'favourites' also makes a lot of sense to have high correlation and is obvious, as well as 'members' and 'scored_by'. So their plots are useless for us now. Let us look at more interesting correlations.

The 'score' and 'chapters' have a good inference that can be drawn from the scatterplot. Mostly, people have given good ratings for Mangas and Books that are having less than 1000 chapters, as we can see that a good number of entries that have good scores are accumulated near the score ranges of 7 to 9 and all those entries' chapter counts are well below 1000. Although, there are a few outliers who are scored considerably well and they have more than 1000 chapters, so the inference is that as the number of chapters increases, people definitely don't like the Manga but that is not all, there are a few more attributes that are also considered like how interesting the story is, for example.




The 'scores' and 'volumes' also will have the same inference. People do not give good scores for Mangas and Books that exceed 50 chapters but still we have outliers which again means that the number of volumes is alone not the criteria. But based on most of


the entries, we can infer that 50 and below is a good number of volumes for a Manga or a Book.




The next scatterplot shows that the score is increasing as the scored_by is also increasing but this is not always an obvious behavior as we know that a product with really good average score that has only very few people who have scored them is not a good judge of the quality of the product. Here, we can see that there are many entries that have good overall scores even after having so many people rating them as those mangas and books are actually nice and people have liked them. Still, this also has some outliers.



The 'score' and 'members' features' correlation makes a good observation. As the manga score is better, more people have added it to their database which means that they have either planned to read it or have already read it. This can mean the other way around too. People who have added it to their database could have also rated it, so the previous plot is a subset of this plot. We can clearly see that the member count increases as the score increases too.



This is also a subset of the previous plot, the number of favorites is increasing as the score is increasing. This is a very obvious relation, because when people like it, they will add it to their favorites list and also rate it if they are interested in doing so. So it perfectly makes sense that the mangas and books with the highest scores also are favorited by more people.



Inferences from the Above Scatter Plots: (We are only concerned with Manga here)

1. **Sazae-san** is the SFW manga with the **highest chapter count** having **6477** chapters in total while **Infection** is the NSFW manga with the highest chapter count having **267** chapters in total.

2. **Jojo's Bizarre Adventures** is the **highest rated** SFW manga with a score of **9.30** and **Sotsugyousei** is the highest rated NSFW manga with a score of **8.51**

3. **Kochira...** is the SFW manga with the **highest volume count** with **200** volumes in total and **Infection** is the NSFW manga with the highest volume count with **30** volumes in total.

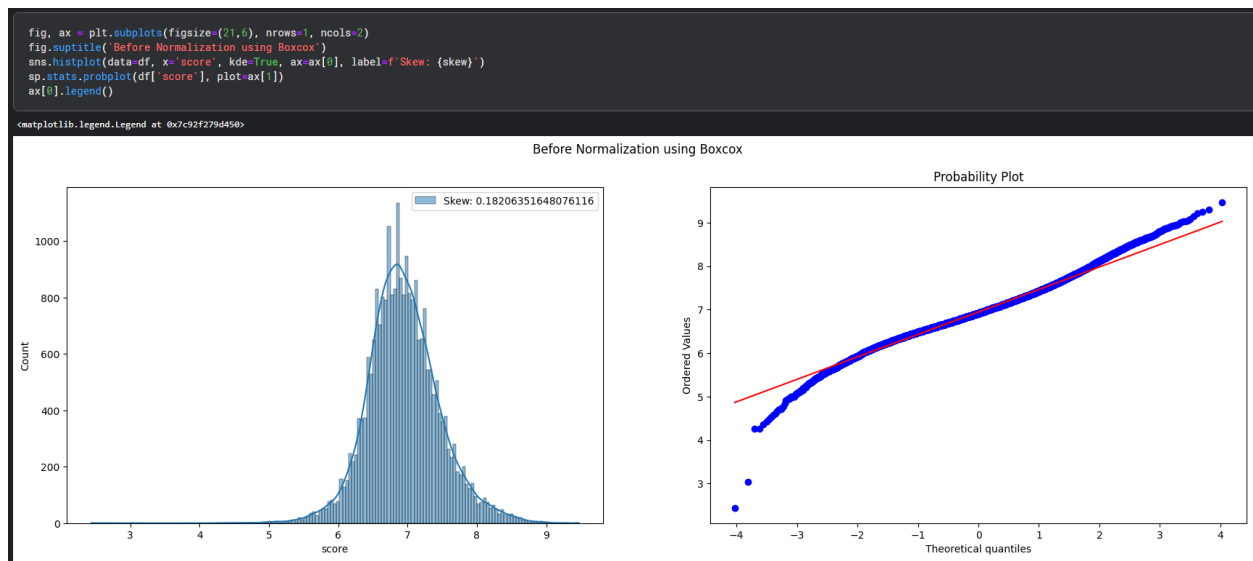
4. **Berserk** (SFW) and **Nozoki Ana** (NSFW) are the manga with the **highest number of people having added them to their database**. They are also the mangas with the **highest favorites count**.

Now, time to normalize the 'score' feature

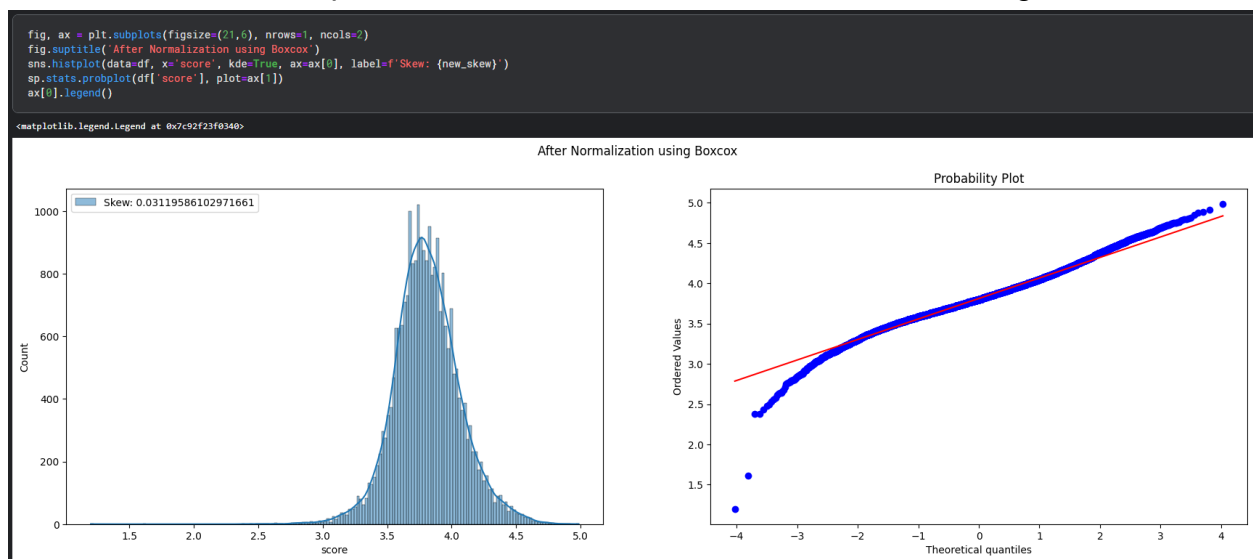
I used the boxcox technique given in the scipy library. The boxcox transform is given as:

$$y = \frac{x^{\lambda-1}}{\lambda}, \lambda \neq 0$$
$$\log(x), \lambda = 0$$

This is how the plot is before the 'score' feature is normalized



This is how the plot is after the 'score' feature is normalized using Boxcox



The data is normalized even further and we can see that the skew is 0.03. But now that we have normalized the score feature, let us check how it change the summary statistics.

	score	scored_by	volumes	chapters	start_date	end_date	members	favorites	real_start_date	real_end_date
count	24636.000000	24636.000000	18322.000000	19939.000000	23765.0	19599.0	24636.000000	24636.000000	23765	19599
mean	3.811176	1769.555691	3.921242	26.740358	2010.219903	2011.090719	4427.903190	143.932619	2010-08-30 10:55:51.718914304	2011-07-14 16:11:40.933721088
min	1.199673	100.000000	1.000000	1.000000	1946.0	1947.0	166.000000	0.000000	1946-01-04 00:00:00	1947-04-01 00:00:00
25%	3.653122	203.000000	1.000000	5.000000	2006.0	2007.0	610.000000	1.000000	2006-09-07 00:00:00	2007-11-28 12:00:00
50%	3.797342	418.000000	2.000000	9.000000	2011.0	2012.0	1191.000000	4.000000	2011-10-05 00:00:00	2012-09-25 00:00:00
75%	3.963639	1060.000000	4.000000	26.000000	2016.0	2017.0	2896.000000	19.000000	2016-10-09 00:00:00	2017-02-07 12:00:00
max	4.986949	394362.000000	200.000000	6477.000000	2023.0	2023.0	643969.000000	119470.000000	2023-07-17 00:00:00	2023-07-28 00:00:00
std	0.256212	8839.552408	6.005744	70.176909	8.462833	8.097007	17656.467096	1711.092022	NaN	NaN

We can see that now the range of scores is not from 0 to 10 anymore, it is 0 to 5 now. Now that we have done these inferences, let us construct a confidence interval for the mean of scores and also let us put forward an hypothesis and test it.

Construction of confidence interval for mean of scores

We see that the population size is large and it follows a normal distribution, we construct a confidence interval using S (since the population standard deviation is unknown). The population mean is given as 3.811176. So we use the formula

$$\bar{x} \pm Z_{\frac{\alpha}{2}} \frac{S}{\sqrt{n}}$$

to construct the 95% confidence interval for the sample mean for the first 100 samples from the population.

Calculating the population mean and the standard deviation S is done as follows:

```
x_bar = df['score'].mean()
S2 = 0
for score in df['score']:
    S2 += (1/df.shape[0]) * (score - x_bar)**2
S = np.sqrt(S2)
print("x_bar: ", x_bar)
print("S: ", S)

x_bar: 3.8111758698644556
S: 0.2562064943153626
```

The confidence interval is constructed as follows:

```
alpha = 0.05
bounds = []
lb = (x_bar - (sp.stats.norm.ppf(1 - alpha/2) * (S / np.sqrt(df.head(100).shape[0]))))
ub = (x_bar + (sp.stats.norm.ppf(1 - alpha/2) * (S / np.sqrt(df.head(100).shape[0]))))
bounds.append(lb)
bounds.append(ub)
```

+ Code

+ Markdown

```
print("The Confidence Interval for Sample Mean is: ", bounds)
```

The Confidence Interval for Sample Mean is: [3.760960319718118, 3.861391420010793]

This is the required 95% confidence interval for the sample mean : [3.760960319718118, 3.861391420010793].

Testing Hypothesis

I propose a hypothesis for the mean of the data ('score' feature)

$$H_0 : \mu = 3.911175 \text{ vs } H_1 : \mu \neq 3.911175$$

With $\alpha = 0.05$, and sigma unknown. So we do a Z-test with S as our population standard deviation as our sample size is 100. Z_{cal} is calculated as

$$Z_{cal} = \frac{\bar{x} - \mu_0}{\frac{S}{\sqrt{n}}}$$

```
Z_cal = (x_bar - 3.9111758698644556)/(S/np.sqrt(100))
Z_cal
```

-3.9031016862871115


```
p_value_z = 2 * (1 - (sp.stats.norm.cdf(np.abs(Z_cal))))  
p_value_z
```

```
9.496777963780012e-05
```

Since the p-value is less than 0.5 (alpha), we reject the null hypothesis.

This concludes the task with the Manga Dataset.